


Durable Incident-Response Agent


Measured comparison of a stateless loop and a checkpointed, approval-gated workflow

Mode: openai | Held-out incidents: 16 | Repetitions: 2 | Seed: 20260802




Interpretation

Diagnosis and remediation scores test task quality. Controlled completion additionally credits explicit policy blocks, operator rejection, and successful compensation instead of treating every non-execution as an undifferentiated failure.



Metric	Baseline	Durable
Executed-crash recovery	100.0% (n=16)	100.0% (n=12)
Cases with duplicate effects	50.0%	0.0%
Failed actions compensated	0.0%	0.0%



Limits

The action tools are simulations, the benchmark is controlled, and deterministic mode validates workflow mechanics rather than general model intelligence. API conclusions require the recorded model configuration and repeated held-out evaluation.