


Governed Text-to-SQL Analyst

Measured one-shot versus governed comparison



Metric	One-shot	Governed
Result accuracy	18.5%	37.0%
Unsafe blocked	0.0%	38.5%
Repair success	80.0%	100.0%

Interpretation

Result hashes measure semantic execution on a fixed synthetic snapshot. Safety results cover the included PII, region, ambiguity, and stored-instruction cases; they are not production certification.