


Long-Term Memory Agent

Measured recent-window, episodic, and hybrid comparison



Metric	Recent window	Episodic	Hybrid
Exact match	0.0%	0.0%	0.0%
Token F1	0.6%	10.0%	6.6%
Evidence recall	1.2%	15.0%	13.1%
Context tokens	90.5	91.1	87.2

Interpretation

Dense episodic retrieval led in this local-model run (10.0% token F1). Adding lexical and recency signals in the hybrid ranker reduced relevance, so its fusion weights need tuning. The run used 150 model calls at an estimated \$0.000.

Deletion compliance

Verified deletion across events, derived facts, retrieval, and tombstones: 100.0%.

Limits

The fixed LoCoMo subset is compact, exact match is strict, and the 7B local answer model limits answer quality even when dense retrieval finds relevant evidence. Model-grounded QA does not establish provider-level deletion guarantees.