


Long-Term Memory Agent

Measured recent-window, episodic, and hybrid comparison



Metric	Recent window	Episodic	Hybrid
Exact match	0.0%	0.0%	0.0%
Token F1	3.2%	11.5%	9.4%
Evidence recall	1.2%	22.3%	19.2%
Context tokens	90.5	71.3	69.6

Interpretation

Episodic retrieval led on this lexical baseline (11.5% token F1). The hybrid recency term reduced relevance, so temporal weighting needs tuning rather than being assumed beneficial. The run used 240 model calls at an estimated \$0.214.

Deletion compliance

Verified deletion across events, derived facts, retrieval, and tombstones: 100.0%.

Limits

The fixed LoCoMo subset is compact, exact match is strict, and lexical retrieval limits semantic recall. Model-grounded QA does not establish provider-level deletion guarantees.